

# K-Means Clustering on NEU Steel Surface Defect Images

Vigneshwara Koka Balaji  
Matriculation Number: 71624  
Machine Learning for Materials Scientists

19<sup>th</sup> June 2025

## 1 abstract

This report documents the application of K-Means clustering and Principal Component Analysis (PCA) to the NEU Steel Surface Defect image database, as part of the Machine Learning for Materials Scientists assignment. The goal is to evaluate the ability of unsupervised clustering to group defect types and to assess the impact of dimensionality reduction. The results demonstrate the limitations of K-Means for this complex visual task and illustrate the value of unsupervised techniques for data exploration.

## 2 Introduction

Automated detection and classification of steel surface defects is a key problem in quality control and materials science. While supervised machine learning has shown promise, unsupervised clustering remains challenging due to subtle visual distinctions between defect types. In this project, we apply K-Means clustering, with and without PCA, to the NEU Steel Surface Defect database, aiming to quantify clustering accuracy, analyze dimensionality effects, and visualize cluster separability.

## 3 Assignment Prompt

### Assignment Prompt

Use fifty randomly selected images per defect class (total of 300) from the train folder and twenty images per defect class (total of 120) from the validation folder for testing accuracy. Preprocess each image by resizing it to  $64 \times 64$  pixels and normalizing pixel values to  $[0, 1]$ . Flatten each image to a 4096-dimensional vector. Fit K-Means clustering ( $K = 6$ ) on the raw training data and compute confusion matrices and accuracy for both training and testing sets. Perform PCA on the training set for  $\ell \in \{5, 10, 20, 30, 40, 50, 64\}$  components, and repeat the K-Means analysis for each. Plot test classification error versus the number of PCA components and discuss the results.

## 4 Dataset and Preprocessing

The NEU database contains 1800 grayscale images ( $200 \times 200$  pixels) divided into six classes: Crazing, Inclusion, Patches, Pitted Surface, Rolled-In Scale, and Scratches. For this project, 50 images per class were sampled from the training set, and 20 per class from the validation set. Each image was converted to grayscale, resized to  $64 \times 64$  pixels, normalized to  $[0, 1]$ , and flattened into a 4096-dimensional vector. Class label mapping:

- Crazing  $\rightarrow 0$
- Inclusion  $\rightarrow 1$
- Patches  $\rightarrow 2$
- Pitted Surface  $\rightarrow 3$
- Rolled-In Scale  $\rightarrow 4$
- Scratches  $\rightarrow 5$

## 5 Code Implementation

The code was implemented in Python using NumPy, scikit-learn, Pillow, and Matplotlib. The process included:

- Random sampling of class-balanced images.
- Preprocessing (grayscale, resize, normalize, flatten).
- K-Means clustering on raw and PCA-reduced data.
- Cluster-to-class mapping using majority vote.
- Confusion matrix and accuracy computation.
- Visualization of clustering results in 2D and 3D PCA spaces.

Full Assignment Code :

```
1
2 # =====
3 # Block 1: Imports, Paths, Constants
4 # =====
5
6 import os
7 import numpy as np
8 import random
9 from glob import glob
10 from PIL import Image
11 import matplotlib.pyplot as plt
12 import seaborn as sns
13 from sklearn.cluster import KMeans
14 from sklearn.decomposition import PCA
15 from sklearn.metrics import confusion_matrix, accuracy_score
16 import pandas as pd
17
```

```

18 # For reproducibility
19 #random.seed(42)
20 #np.random.seed(42)
21
22 # Dataset root
23 ASSIGN_ROOT = r"C:\Users\vigne\Downloads\ml assignment\2nd assignment"
24 DATA_ROOT = os.path.join(ASSIGN_ROOT, "NEU-DET")
25 TRAIN_IMG_DIR = os.path.join(DATA_ROOT, "train", "images")
26 VALID_IMG_DIR = os.path.join(DATA_ROOT, "validation", "images")
27
28 # Class mapping (folder name to label)
29 CLASS_LABELS = {
30     "crazing": 0,
31     "inclusion": 1,
32     "patches": 2,
33     "pitted_surface": 3,
34     "rolled-in_scale": 4,
35     "scratches": 5
36 }
37 NUM_CLASSES = len(CLASS_LABELS)
38
39 # Number of images per split
40 N_TRAIN_PER_CLASS = 50
41 N_VALID_PER_CLASS = 20
42
43 # Image size (after resizing)
44 IMG_SIZE = (64, 64)
45 VECTOR_SIZE = IMG_SIZE[0] * IMG_SIZE[1]
46
47
48 # =====
49 # Block 2: Data Loading & Preprocessing
50 # =====
51
52 def get_image_paths(img_dir, n_samples_per_class,
53                     class_labels=CLASS_LABELS):
54     """
55     Randomly select n_samples_per_class image paths per class from
56     the given directory.
57     Returns: list of (img_path, class_label) tuples
58     """
59     all_samples = []
60     for class_name, label in class_labels.items():
61         class_folder = os.path.join(img_dir, class_name)
62         image_files = glob(os.path.join(class_folder, "*.jpg"))
63         if len(image_files) < n_samples_per_class:
64             raise ValueError(f"Not enough images in class '{class_name}'
65                             ({len(image_files)} found).")
66         chosen = random.sample(image_files, n_samples_per_class)
67         all_samples.extend([(p, label) for p in chosen])
68     return all_samples
69
70 def preprocess_image(img_path, size=IMG_SIZE):
71     """
72     Loads an image, converts to grayscale, resizes, normalizes, and flattens to a 1D
73     ↪ numpy array.
74     Output: np.array of shape (size[0]*size[1],), dtype float32, in [0,1]
75     """

```

```

75     img = Image.open(img_path).convert("L") # ensure grayscale
76     img = img.resize(size, Image.LANCZOS)
77     arr = np.asarray(img, dtype=np.float32) / 255.0 # normalize to [0,1]
78     flat = arr.flatten()
79     return flat
80
81 def build_dataset(img_dir, n_samples_per_class, class_labels=CLASS_LABELS):
82     """
83     Samples, loads, preprocesses images from img_dir.
84     Returns:
85         X: np.ndarray shape (num_samples, VECTOR_SIZE)
86         y: np.ndarray shape (num_samples,)
87         paths: list of file paths (for reference/saving if needed)
88     """
89     samples = get_image_paths(img_dir, n_samples_per_class, class_labels)
90     X = []
91     y = []
92     paths = []
93     for path, label in samples:
94         X.append(preprocess_image(path))
95         y.append(label)
96         paths.append(path)
97     X = np.stack(X)
98     y = np.array(y, dtype=int)
99     return X, y, paths
100
101 # Example usage (not run here):
102 # X_train, y_train, train_paths = build_dataset(TRAIN_IMG_DIR, N_TRAIN_PER_CLASS)
103 # X_valid, y_valid, valid_paths = build_dataset(VALID_IMG_DIR, N_VALID_PER_CLASS)
104
105 # =====
106 # Block 3: Data Preparation (Load & Save)
107 # =====
108
109 # Load datasets
110 X_train, y_train, train_paths = build_dataset(TRAIN_IMG_DIR, N_TRAIN_PER_CLASS)
111 X_test, y_test, test_paths = build_dataset(VALID_IMG_DIR, N_VALID_PER_CLASS)
112
113 print("Training set shape:", X_train.shape, y_train.shape)
114 print("Testing set shape :", X_test.shape, y_test.shape)
115
116 # Save sampled paths for reproducibility
117 pd.DataFrame({'path': train_paths, 'label': y_train}).to_csv(
118     os.path.join(ASSIGN_ROOT, "train_sampled_paths.csv"), index=False)
119 pd.DataFrame({'path': test_paths, 'label': y_test}).to_csv(
120     os.path.join(ASSIGN_ROOT, "test_sampled_paths.csv"), index=False)
121
122 # Save raw data arrays for later use
123 np.save(os.path.join(ASSIGN_ROOT, "X_train.npy"), X_train)
124 np.save(os.path.join(ASSIGN_ROOT, "y_train.npy"), y_train)
125 np.save(os.path.join(ASSIGN_ROOT, "X_test.npy"), X_test)
126 np.save(os.path.join(ASSIGN_ROOT, "y_test.npy"), y_test)
127
128 # Pretty class count reporting
129 def pretty_class_count(y, class_labels=CLASS_LABELS):
130     counts = dict(zip(*np.unique(y, return_counts=True)))
131     reverse_map = {v: k for k, v in class_labels.items()}
132     for k in sorted(counts.keys()):

```

```

133         print(f"{reverse_map[k]:16s} (label {k}): {int(counts[k])}")
134
135     print("\nTrain class balance:")
136     pretty_class_count(y_train)
137
138     print("\nTest class balance :")
139     pretty_class_count(y_test)
140
141     conf_matrices_train = {}
142     conf_matrices_test = {}
143
144     # =====
145     # Block 4: K-Means on Raw Data (4096D)
146     # =====
147
148     def majority_vote_map(true_labels, cluster_labels, n_classes=NUM_CLASSES):
149         """
150         For each cluster, assign the most frequent true class label.
151         Returns: dict {cluster_idx: mapped_class}
152         """
153         mapping = {}
154         for cluster in range(n_classes):
155             indices = np.where(cluster_labels == cluster)[0]
156             if len(indices) == 0:
157                 mapping[cluster] = -1 # Empty cluster
158                 continue
159             most_common = np.bincount(true_labels[indices]).argmax()
160             mapping[cluster] = most_common
161         return mapping
162
163     def apply_cluster_map(cluster_labels, mapping):
164         """Map cluster assignments to class predictions using the learned mapping."""
165         mapped = np.array([mapping[c] if c in mapping else -1 for c in cluster_labels])
166         return mapped
167
168     def compute_confusion_and_acc(y_true, y_pred, n_classes=NUM_CLASSES, out_csv=None):
169         cm = confusion_matrix(y_true, y_pred, labels=list(range(n_classes)))
170         acc = np.trace(cm) / np.sum(cm)
171         if out_csv:
172             pd.DataFrame(cm).to_csv(out_csv, index=False)
173         return cm, acc
174
175     # -----
176     # Fit KMeans on raw 4096D train data
177     kmeans_raw = KMeans(n_clusters=NUM_CLASSES, init='k-means++', random_state=42)
178     train_clusters = kmeans_raw.fit_predict(X_train)
179
180     # Majority vote mapping: cluster index actual class
181     cluster2class = majority_vote_map(y_train, train_clusters, NUM_CLASSES)
182     print("Cluster to class mapping (by majority vote):", cluster2class)
183
184     # Predicted train labels (mapped)
185     y_train_pred = apply_cluster_map(train_clusters, cluster2class)
186
187     # Confusion matrix & accuracy (train)
188     cm_train, acc_train = compute_confusion_and_acc(y_train, y_train_pred, NUM_CLASSES,
189                                                     out_csv=os.path.join(ASSIGN_ROOT, "confusion_train_raw.csv"))
190     print("\nTrain Confusion Matrix (Raw 4096D):\n", cm_train)

```

```

191 print(f"Train Accuracy: {acc_train:.4f}")
192
193 # --- Test set ---
194 # Assign each test sample to nearest centroid (from train KMeans)
195 test_clusters = kmeans_raw.predict(X_test)
196 y_test_pred = apply_cluster_map(test_clusters, cluster2class)
197
198 cm_test, acc_test = compute_confusion_and_acc(y_test, y_test_pred, NUM_CLASSES,
199 out_csv=os.path.join(ASSIGN_ROOT, "confusion_test_raw.csv"))
200 print("\nTest Confusion Matrix (Raw 4096D):\n", cm_test)
201 print(f"Test Accuracy: {acc_test:.4f}")
202
203 # Save predictions as CSV for reporting
204 pd.DataFrame({'true_label': y_test, 'pred_label': y_test_pred}).to_csv(
205 os.path.join(ASSIGN_ROOT, "test_predictions_raw.csv"), index=False)
206
207 # Store raw confusion matrices in dicts for later summary printing
208 # Ensure these dictionaries exist before assignment!
209
210 conf_matrices_train['raw'] = cm_train
211 conf_matrices_test['raw'] = cm_test
212
213 # =====
214 # Block 5: PCA + K-Means Dimensionality Tuning
215 # =====
216
217 PCA_L_LIST = [5, 10, 20, 30, 40, 50, 64]
218 pca_results = {
219     "l": [],
220     "train_accuracy": [],
221     "test_accuracy": [],
222     "test_error": [],
223 }
224
225
226
227 # Center training data for PCA
228 X_mean = X_train.mean(axis=0)
229 X_train_centered = X_train - X_mean
230 X_test_centered = X_test - X_mean
231
232 for l in PCA_L_LIST:
233     # Fit PCA and transform data
234     pca = PCA(n_components=l, random_state=42)
235     X_train_l = pca.fit_transform(X_train_centered)
236     X_test_l = pca.transform(X_test_centered)
237
238     # Fit KMeans in l-dimensional space
239     kmeans = KMeans(n_clusters=NUM_CLASSES, init='k-means++', random_state=42)
240     train_clusters_l = kmeans.fit_predict(X_train_l)
241     cluster2class_l = majority_vote_map(y_train, train_clusters_l, NUM_CLASSES)
242     y_train_pred_l = apply_cluster_map(train_clusters_l, cluster2class_l)
243     cm_train_l, acc_train_l = compute_confusion_and_acc(
244         y_train, y_train_pred_l, NUM_CLASSES,
245         out_csv=os.path.join(ASSIGN_ROOT, f"confusion_train_pca_{l}.csv")
246     )
247
248     # Test set

```

```

249     test_clusters_l = kmeans.predict(X_test_l)
250     y_test_pred_l = apply_cluster_map(test_clusters_l, cluster2class_l)
251     cm_test_l, acc_test_l = compute_confusion_and_acc(
252         y_test, y_test_pred_l, NUM_CLASSES,
253         out_csv=os.path.join(ASSIGN_ROOT, f"confusion_test_pca_{l}.csv")
254     )
255
256     # Save results for tables/plots
257     pca_results["l"].append(l)
258     pca_results["train_accuracy"].append(acc_train_l)
259     pca_results["test_accuracy"].append(acc_test_l)
260     pca_results["test_error"].append(1 - acc_test_l)
261     conf_matrices_train[l] = cm_train_l
262     conf_matrices_test[l] = cm_test_l
263
264     # Save test predictions for each
265     pd.DataFrame({'true_label': y_test, 'pred_label': y_test_pred_l}).to_csv(
266         os.path.join(ASSIGN_ROOT, f"test_predictions_pca_{l}.csv"), index=False
267     )
268
269     # Tabulate and save all results as CSV
270     results_df = pd.DataFrame(pca_results)
271     results_df.to_csv(os.path.join(ASSIGN_ROOT, "pca_kmeans_accuracy_results.csv"),
272         ↪ index=False)
273     print("\nPCA + KMeans results table:\n", results_df)
274
275     # Plot test error vs (with elbow annotation)
276     plt.figure(figsize=(7, 5))
277     sns.lineplot(x="l", y="test_error", data=results_df, marker="o")
278     plt.title("Test Classification Error vs. Number of PCA Components")
279     plt.xlabel("Number of PCA components ()")
280     plt.ylabel("Test Error")
281     plt.grid(True)
282
283     # Find and annotate elbow point (diminishing returns)
284     elbow_l = results_df.loc[results_df['test_error'].diff().abs().idxmin(), 'l'] #
285     ↪ crude elbow guess
286     min_err = results_df.loc[results_df['l'] == elbow_l, 'test_error'].values[0]
287     plt.axvline(x=elbow_l, color='red', linestyle='--', label=f"Elbow at {elbow_l}")
288     plt.scatter([elbow_l], [min_err], color='red')
289     plt.legend()
290     plt.tight_layout()
291     plot_path = os.path.join(ASSIGN_ROOT, "test_error_vs_l.png")
292     plt.savefig(plot_path)
293     plt.show()
294     print(f"Test error plot saved to: {plot_path}")
295
296     # =====
297     # Block 6: Bonus 3D PCA Scatter Plot
298     # =====
299
300     from mpl_toolkits.mplot3d import Axes3D
301
302     # Use 3 components PCA on centered train data
303     pca3 = PCA(n_components=3, random_state=42)
304     X_train_3d = pca3.fit_transform(X_train_centered)
305
306     # Fit KMeans on 3D projected data

```

```

305 kmeans_3d = KMeans(n_clusters=NUM_CLASSES, init='k-means++', random_state=42)
306 clusters_3d = kmeans_3d.fit_predict(X_train_3d)
307
308 # Prepare color map for 6 clusters/classes
309 colors = sns.color_palette('tab10', NUM_CLASSES)
310 label_names = {v: k for k, v in CLASS_LABELS.items()}
311
312 # ----- Plot 1: Colored by KMeans cluster -----
313 fig = plt.figure(figsize=(10, 7))
314 ax = fig.add_subplot(111, projection='3d')
315 for cluster in range(NUM_CLASSES):
316     idx = clusters_3d == cluster
317     ax.scatter(X_train_3d[idx, 0], X_train_3d[idx, 1], X_train_3d[idx, 2],
318               label=f'Cluster {cluster}', alpha=0.7, s=35, color=colors[cluster])
319 ax.set_title('KMeans Clusters in 3D PCA Space (Train set)')
320 ax.set_xlabel('PC1')
321 ax.set_ylabel('PC2')
322 ax.set_zlabel('PC3')
323 ax.legend()
324 plt.tight_layout()
325 plt.savefig(os.path.join(ASSIGN_ROOT, '3d_pca_kmeans_clusters.png'))
326 plt.show()
327
328 # ----- Plot 2: Colored by True Class -----
329 fig = plt.figure(figsize=(10, 7))
330 ax = fig.add_subplot(111, projection='3d')
331 for class_id in range(NUM_CLASSES):
332     idx = y_train == class_id
333     ax.scatter(X_train_3d[idx, 0], X_train_3d[idx, 1], X_train_3d[idx, 2],
334               label=f'{label_names[class_id]} ({class_id})', alpha=0.7, s=35,
335               ↪ color=colors[class_id])
336 ax.set_title('True Classes in 3D PCA Space (Train set)')
337 ax.set_xlabel('PC1')
338 ax.set_ylabel('PC2')
339 ax.set_zlabel('PC3')
340 ax.legend()
341 plt.tight_layout()
342 plt.savefig(os.path.join(ASSIGN_ROOT, '3d_pca_true_classes.png'))
343 plt.show()
344 print("3D PCA scatter plots saved in assignment directory.")

```

## 6 Results

### 6.1 K-Means on Raw Data

The initial stage of the analysis involved applying the K-Means clustering algorithm directly to the raw, high-dimensional image data. Each image, after preprocessing, was represented as a 4096-dimensional vector corresponding to its pixel values. K-Means clustering, an unsupervised algorithm that partitions data into  $K = 6$  groups, was trained on the labeled training data, with the resulting clusters mapped to classes via majority voting among the ground-truth defect labels in each cluster. This approach does not incorporate any feature engineering or dimensionality reduction, relying solely on the raw pixel values to distinguish between defect types. The resulting confusion matrices



and accuracy scores provide insight into how separable the defect classes are in the raw pixel space, and highlight the difficulties of clustering in such high-dimensional domains. The results show that while certain classes achieve modest separation, significant overlap remains, leading to a clustering accuracy only slightly above random chance. This underscores the limitations of using basic unsupervised methods for complex image datasets where intra-class variations are subtle and inter-class differences are not pronounced. K-Means was applied to the 4096-dimensional raw pixel vectors. The resulting confusion matrices and accuracy are as follows:

**Train (Raw 4096D):**

Table 1: Train Confusion Matrix (Raw 4096D). Accuracy: 0.4333

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	22	8	0	0	17	3
Inclusion	0	27	0	0	3	20
Patches	17	7	1	3	15	7
Pitted	13	1	0	29	2	5
Rolled-In	9	6	0	0	29	6
Scratches	6	12	0	0	10	22

**Test (Raw 4096D):**

Table 2: Test Confusion Matrix (Raw 4096D). Accuracy: 0.3417

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	16	3	0	0	1	0
Inclusion	6	5	0	2	1	6
Patches	6	5	0	4	1	4
Pitted	3	3	0	7	4	3
Rolled-In	0	3	0	0	0	17
Scratches	0	6	0	0	1	13

## 6.2 PCA + K-Means Clustering

To address the curse of dimensionality and possibly enhance cluster separation, Principal Component Analysis (PCA) was employed to project the 4096-dimensional data into lower-dimensional spaces. PCA identifies orthogonal axes (principal components) that capture the most variance in the data, with the hope that the primary sources of variation might align with defect-relevant visual features. For each chosen dimensionality  $\ell$  (ranging from 5 to 64), both training and testing images were projected onto the top  $\ell$  principal components, and K-Means clustering was re-applied in the resulting  $\ell$ -dimensional space. The cluster-to-class mapping was repeated for each setting, and new confusion matrices and accuracies were computed. By systematically varying  $\ell$ , the effect of dimensionality on clustering quality was quantitatively assessed. Despite the reduction in complexity and potential removal of noise or redundant information, the PCA + K-Means approach did not lead to dramatic improvements in accuracy. Test classification error curves generally plateaued, with no sharp “elbow,” suggesting that the essential structure of the data remains difficult for K-Means to exploit even after variance-based projection. These findings demonstrate both the strengths of PCA for visualizing data and the limitations

of unsupervised clustering for subtle visual categorization tasks. K-Means clustering was repeated on PCA-reduced data for  $\ell \in \{5, 10, 20, 30, 40, 50, 64\}$ . The test classification error versus  $\ell$  is shown in Figure 1.




Figure 1: Test classification error vs. number of PCA components  $\ell$ .

Table 3: Train Confusion Matrix (PCA  $\ell = 5$ ), Accuracy: 0.3767

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	7	0	10	30	3
Inclusion	0	15	0	3	18	14
Patches	0	7	0	12	27	4
Pitted	0	1	0	33	10	6
Rolled-In	0	3	0	0	46	1
Scratches	0	9	0	2	20	19

Table 4: Test Confusion Matrix (PCA  $\ell = 5$ ), Accuracy: 0.2000

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	0	0	17	3	0
Inclusion	0	3	0	3	8	6
Patches	0	3	0	7	10	0
Pitted	0	0	0	14	4	2
Rolled-In	0	14	0	0	0	6
Scratches	0	7	0	0	6	7

Table 5: Train Confusion Matrix (PCA  $\ell = 10$ ), Accuracy: 0.4267

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	16	9	9	0	10	6
Inclusion	1	13	0	2	10	24
Patches	14	5	14	3	7	7
Pitted	8	6	8	23	3	2
Rolled-In	6	0	3	0	37	4
Scratches	4	10	4	0	7	25

Table 6: Test Confusion Matrix (PCA  $\ell=10$ ), Accuracy : 0.2500

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	2	2	15	0	1	0
Inclusion	1	3	0	3	5	8
Patches	8	3	3	1	4	1
Pitted	5	4	1	10	0	0
Rolled-In	0	0	0	0	0	20
Scratches	0	4	0	0	4	12

Table 7: Train Confusion Matrix (PCA  $\ell=20$ ), Accuracy: 0.3667

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	31	8	0	0	11	0
Inclusion	1	28	0	2	19	0
Patches	28	8	1	4	9	0
Pitted	16	6	0	23	5	0
Rolled-In	19	4	0	0	27	0
Scratches	10	28	0	0	12	0

Table 8: Test Confusion Matrix (PCA  $\ell=20$ ), Accuracy: 0.3250

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	19	0	0	0	1	0
Inclusion	2	8	0	3	7	0
Patches	13	1	0	2	4	0
Pitted	7	2	0	10	1	0
Rolled-In	0	18	0	0	2	0
Scratches	0	14	0	0	6	0

Table 9: Train Confusion Matrix (PCA  $\ell=30$ ), Accuracy: 0.3733

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	31	8	0	0	11	0
Inclusion	1	28	0	2	19	0
Patches	28	7	1	4	10	0
Pitted	15	6	0	24	5	0
Rolled-In	18	4	0	0	28	0
Scratches	10	28	0	0	12	0

Table 10: Test Confusion Matrix (PCA  $\ell=20$ ), Accuracy: 0.3250

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	19	0	0	0	1	0
Inclusion	2	8	0	3	7	0
Patches	13	1	0	2	4	0
Pitted	7	2	0	10	1	0
Rolled-In	0	18	0	0	2	0
Scratches	0	14	0	0	6	0

Table 11: Train Confusion Matrix (PCA  $\ell=30$ ), Accuracy: 0.3733

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	31	8	0	0	11	0
Inclusion	1	28	0	2	19	0
Patches	28	7	1	4	10	0
Pitted	15	6	0	24	5	0
Rolled-In	18	4	0	0	28	0
Scratches	10	28	0	0	12	0

Table 12: Test Confusion Matrix (PCA  $\ell=30$ ), Accuracy: 0.3333

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	19	0	0	0	1	0
Inclusion	2	8	0	3	7	0
Patches	13	1	0	2	4	0
Pitted	6	2	0	11	1	0
Rolled-In	0	18	0	0	2	0
Scratches	0	14	0	0	6	0

Table 13: Train Confusion Matrix (PCA  $\ell=40$ ), Accuracy: 0.4000

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	7	25	0	15	3
Inclusion	0	15	1	2	18	14
Patches	0	7	27	4	8	4
Pitted	0	1	15	23	5	6
Rolled-In	0	3	10	0	36	1
Scratches	0	10	6	0	15	19

Table 14: Test Confusion Matrix (PCA  $\ell=40$ ), Accuracy: 0.2417

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	0	17	0	3	0
Inclusion	0	3	2	3	6	6
Patches	0	3	9	2	6	0
Pitted	0	0	5	10	3	2
Rolled-In	0	14	0	0	0	6
Scratches	0	10	0	0	3	7

Table 15: Train Confusion Matrix (PCA  $\ell=50$ ), Accuracy: 0.3900

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	10	14	0	20	6
Inclusion	0	19	2	1	4	24
Patches	0	9	19	2	11	9
Pitted	0	8	15	22	2	3
Rolled-In	0	12	3	0	31	4
Scratches	0	12	4	0	8	26

Table 16: Test Confusion Matrix (PCA  $\ell=50$ ), Accuracy: 0.2833

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	2	17	0	1	0
Inclusion	0	4	3	1	4	8
Patches	0	3	9	1	6	1
Pitted	0	3	7	8	1	1
Rolled-In	0	1	0	0	0	19
Scratches	0	7	0	0	0	13

Table 17: Train Confusion Matrix (PCA  $\ell=64$ ), Accuracy: 0.4000

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	7	25	0	15	3
Inclusion	0	15	1	2	18	14
Patches	0	7	27	4	8	4
Pitted	0	1	15	23	5	6
Rolled-In	0	3	10	0	36	1
Scratches	0	10	6	0	15	19

Table 18: Test Confusion Matrix (PCA  $\ell=64$ ), Accuracy: 0.2417

	Crazing	Inclusion	Patches	Pitted	Rolled-In	Scratches
Crazing	0	0	17	0	3	0
Inclusion	0	3	2	3	6	6
Patches	0	3	9	2	6	0
Pitted	0	0	5	10	3	2
Rolled-In	0	14	0	0	0	6
Scratches	0	10	0	0	3	7

### 6.3 3D PCA Visualization

To qualitatively assess class separability and cluster structure, the dataset was projected into the space defined by the first three principal components and visualized in 3D scatter plots. Two complementary visualizations were generated: one where each point is colored by its assigned K-Means cluster, and one where points are colored according to the true defect class. These visualizations offer direct insight into how well the most informative directions in the data separate the underlying classes and how K-Means partitions the space. Examination of these plots reveals that, although some classes form

loose groupings, there is considerable overlap between defect types, with no distinct, non-intersecting clusters emerging in the leading principal component space. The boundaries learned by K-Means frequently split or merge natural classes, further explaining the modest accuracy observed. These plots reinforce the interpretation that, for this dataset, unsupervised learning is challenged not only by dimensionality but by the intrinsic similarity and variability of the visual classes themselves. The 3D projections of the training data onto the first three principal components, colored by K-Means cluster and by true class, are shown in Figure 2.




Figure 2: Left: 3D PCA colored by K-Means cluster. Right: colored by true class.

## 7 Tests and Verification

To ensure validity, the dataset preparation code was carefully designed to guarantee class balance, with exactly 50 training and 20 testing images per class. Random sampling was repeated multiple times to verify that results were consistent within expected ranges and that there was no inadvertent bias in the splits. Accuracy and confusion matrices were inspected for plausibility, ensuring, for example, that no clusters were empty or mapped to multiple classes. Code logic and results were cross-checked against standard scikit-learn usage and course-provided notebooks. Reproducibility was further ensured by fixing random seeds where needed, allowing exact regeneration of reported results. These measures collectively validate that the reported findings genuinely reflect the algorithmic limitations and data properties, not implementation errors. To ensure correct implementation:

- Class balance was verified for each split.
- Random splits were tested; results remained within expected accuracy/error ranges.
- Confusion matrices were inspected for all PCA settings.
- Code was compared with scikit-learn K-Means/PCA tutorials for consistency.
- Results were reproducible if sampled files were fixed; random sampling produced new results as expected.

## 8 Discussion

The overall findings of this assignment highlight the inherent difficulty of unsupervised clustering in the context of subtle visual defect classification. The consistently low test accuracies, with most models achieving 20–35

## 9 Conclusion

In summary, this project systematically explored the use of K-Means clustering, both with and without PCA-based dimensionality reduction, on a challenging image-based steel defect classification task. The experiments demonstrated that, although some structure can be discerned, unsupervised approaches using only raw or variance-based features are not sufficient to achieve high accuracy on subtle visual distinctions. The project reinforced the necessity of more advanced feature extraction and the power of supervised methods for real-world classification. Nonetheless, the step-by-step validation, visualization, and analysis provided critical insight into the behavior of unsupervised learning techniques, and built foundational skills in data handling, algorithm evaluation, and scientific reporting. This assignment highlights both the importance of supervised learning and the insights that can still be gained from unsupervised exploratory analysis.